Rubric evidence audit - not a grade

All ten categories total 100 marks. This maps evidence to requirements; it does not award marks or certify student authorship.

Category	Marks	Evidence / outstanding
Identity and reproducibility	8	Confirmed name/email; config.py, fingerprint and two-run equality.
Data generation and validation	12	generate_data.py, data_quality.json and validation tests.
Preprocessing and leakage	10	Independent pipelines; stored medians tested against train rows.
Regression	12	Three candidates, metrics, residuals; student error wording pending.
Classification	15	Three candidates, metrics, probabilities and confusion matrices; personal error wording pending.
Clustering	10	K=2..6, silhouette, profiles, labeled PCA; personal interpretation pending.
CLI and FastAPI	10	Saved models, shared schema, actual TestClient and CLI evidence.
Testing and code quality	8	Modular sources; actual pytest and reproducibility captures.
Report and responsible AI	10	Seven-page report, model card and AI disclosure; student reflection pending.
Submission evidence and defense	5	ZIP, Git hash, test log and checklist; video/signature/student defense pending.